
RLLTE: Long-Term Evolution Project of Reinforcement Learning

¹Mingqi Yuan, ²Zequn Zhang, ³Yang Xu, ⁴Shihao Luo, ¹Bo Li, ²Xin Jin, and
²Wenjun Zeng

¹The Hong Kong Polytechnic University

²Eastern Institute for Advanced Study

³Purdue University

⁴Da-Jiang Innovations (DJI)

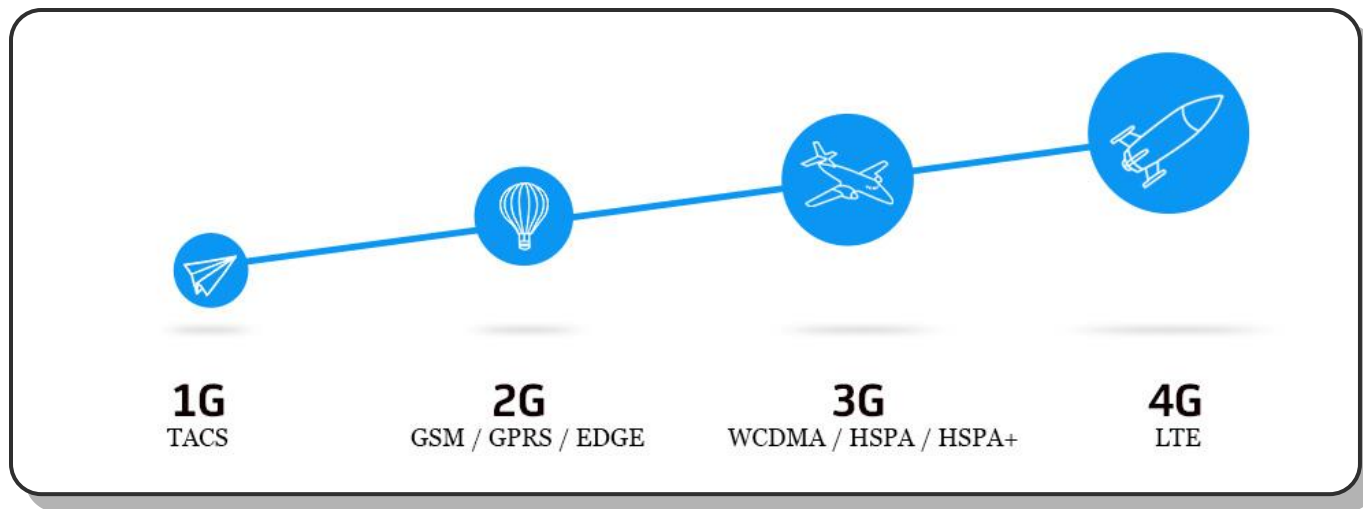
- ❑ Volatile performance of different implementations;
- ❑ Algorithm updates are very complex and miscellaneous;
- ❑ Unfriendly support for the latest tricks;
- ❑ Incomplete benchmark testing;
- ❑ Expensive computational cost of algorithm reproduction;
- ❑ Few active repositories;
- ❑ High learning costs for developers.



What is RLLTE?



- ❑ A novel **reinforcement learning** (RL) library inspired by the **long-term evolution** (LTE) standard project in telecommunications.



- ❑ GitHub Link: <https://github.com/RLE-Foundation/rllte>



What is RLLTE for?













- ❑ Setting common **standards** for RL engineering practice;
- ❑ Accelerating RL algorithms **iteration**;
- ❑ Tracking the **latest** research **progress**;
- ❑ Providing reusable and reliable **baselines**;
- ❑ Achieving the goal of “**RL For Everyone.**”



Highlight Features




- ❑  Large language model-empowered **copilot**;
- ❑  Latest **algorithms** and **tricks**;
- ❑  Standard and sophisticated modules for **redevelopment**;
- ❑  Highly modularized design for complete **decoupling** of RL algorithms;
- ❑  Optimized workflow for full hardware **acceleration**;
- ❑  Support for **custom environments** and **modules**;
- ❑  Support for multiple computing devices like **GPU** and **NPU**;
- ❑  Support for RL model engineering **deployment** (TensorRT, CANN, ...);
- ❑  Large number of reusable **benchmarks** (See [rlte-hub](https://github.com/rlte-hub));



Comparison



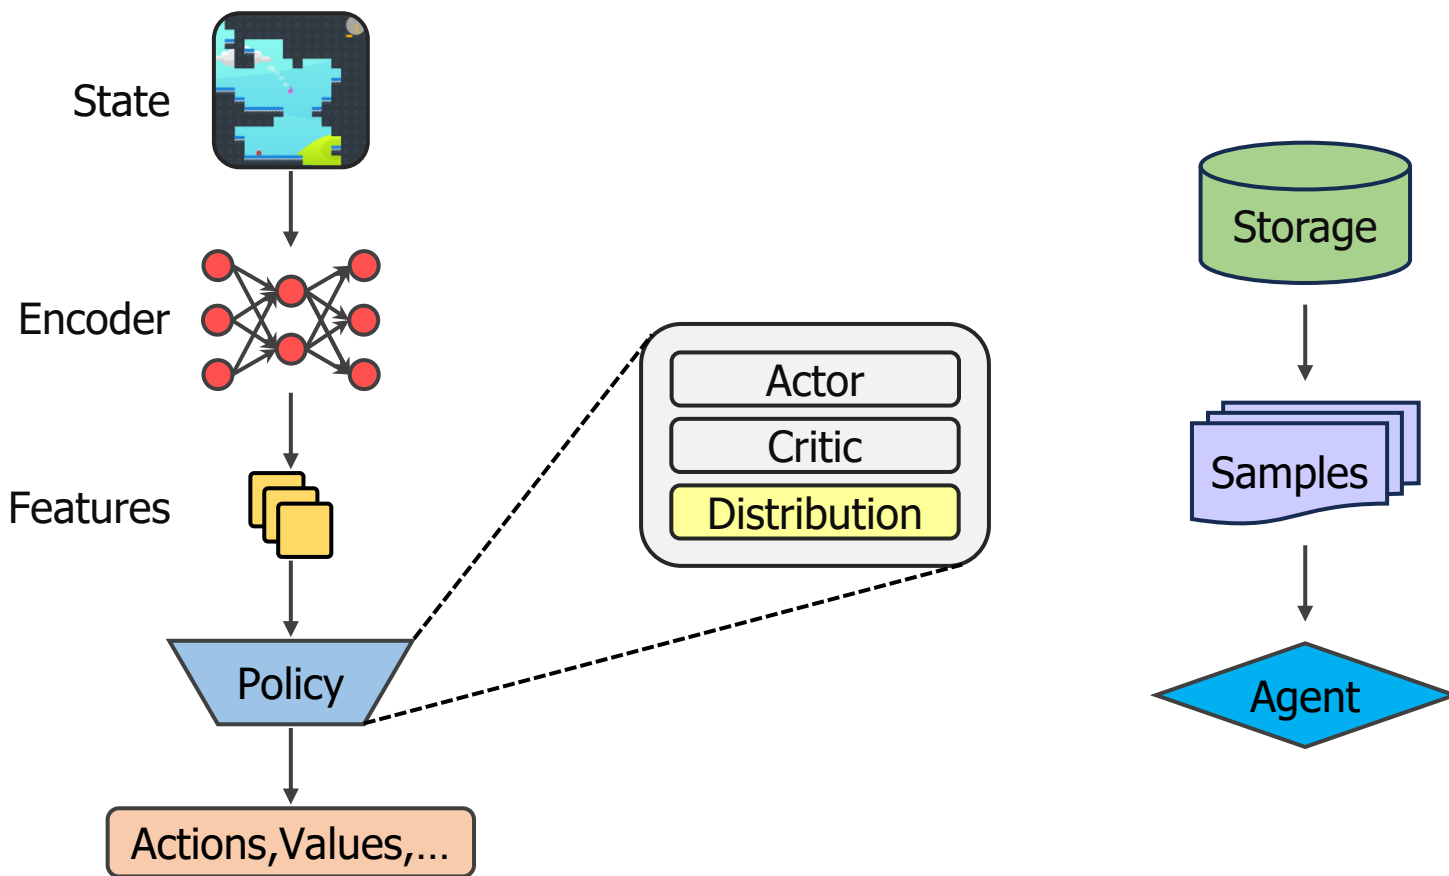
	Modularize d	Custom Env	Gymnasiu m	Custom Module	Data Augmentat ion	Benchmark	Deploymen t	Evaluation	Multi- device	Active
Baselines	✓	✓	✗	—	✗	—	✗	✗	✗	✗
Stable- Baselines 3	✓	✓	✓	—	✗	—	✗	✗	✗	✓
CleanRL	✗	✗	—	✗	✗	—	✗	✗	✗	✓
Ray/rllib	✓	✓	✗	—	✗	—	✗	✗	✗	—
rlypt	✓	✗	✗	✗	✗	—	✗	✗	✗	✗
Tianshou	✓	✓	✓	✗	✗	—	✗	✗	✗	—
ElegantRL	✓	✓	✗	✗	✗	—	✗	✗	✗	—
SpinningU p	✗	✓	✗	✗	✗	—	✗	✗	✗	✗
ACME	✗	✓	✗	✗	✗	—	✗	✗	✗	✗
Dopamine	✗	✗	✗	✗	✗	—	✗	✗	✗	✗
RLLTE	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓



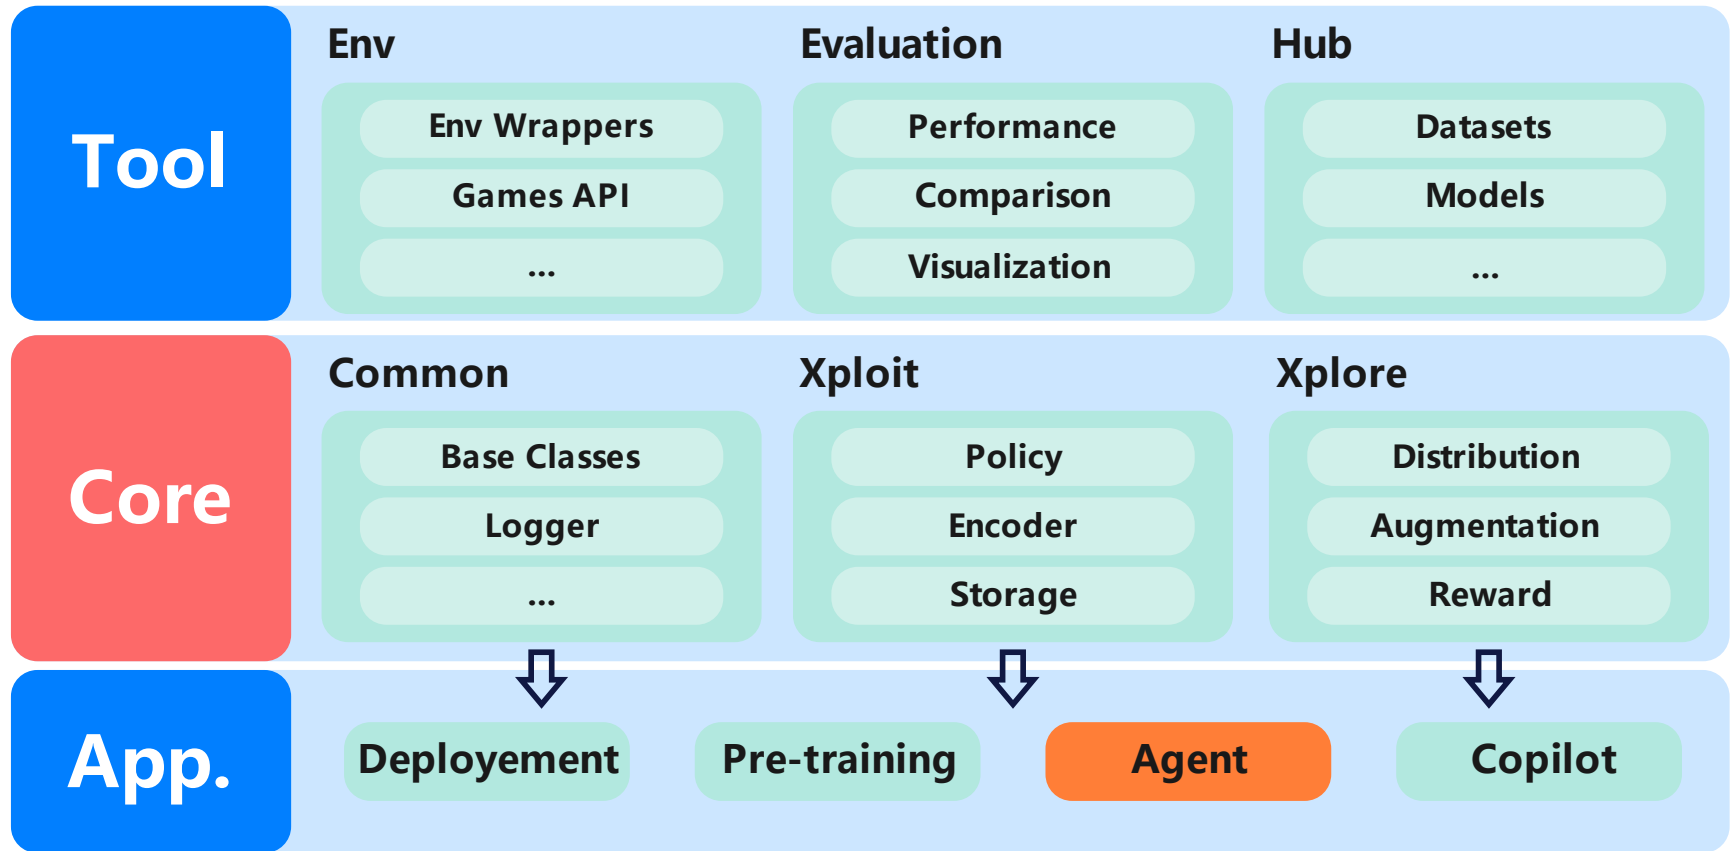
Architecture (Decoupling)



❑ RL Algorithms **Decoupling**



Architecture (Overview)



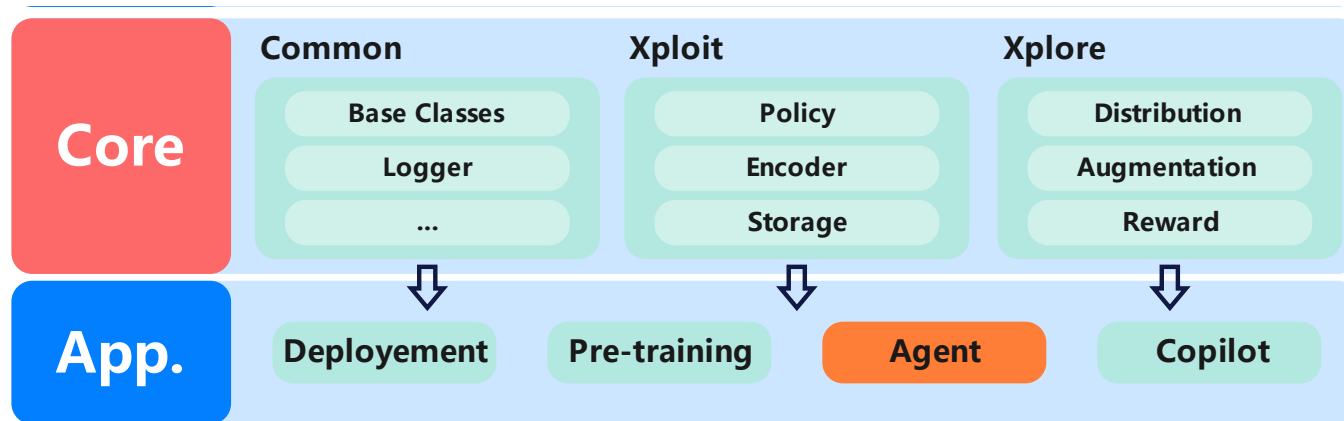
- ❑ **Common:** Base classes and auxiliary modules.
- ❑ **Xploit:** Modules that focus on **exploitation** in RL.
 - **Encoder:** **feature** extraction;
 - **Policy:** **interaction** and learning;
 - **Storage:** experience storage and sampling.
- ❑ **Xplore:** Modules that focus on **exploration** in RL.
 - **Distribution:** **action** sampling;
 - **Augmentation:** observation data augmentation;
 - **Reward:** **intrinsic reward** modules.



Architecture (Application)



- ❑ **Agent:** Implemented RL Agents using **rlite** building blocks.
- ❑ **Pre-Training:** Methods of **pre-training** in RL.
- ❑ **Deployment:** Methods of model **deployment** in RL.
- ❑ **Copilot:** **LLM-based copilot** that helps developer build RL applications with **rlite**.



Architecture (Application~Agent)



- ❑ Modules-Oriented ---> Algorithms-Oriented;
- ❑ RL algorithms are the **applications** of basic modules.

```
from rllte.xploit.encoder import MnihCnnEncoder
```

```
from rllte.xploit.storage import VanillaRolloutStorage
```

```
from rllte.xploit.policy import OnPolicySharedActorCritic
```

```
from rllte.xplore.distribution import Categorical
```

```
def update(self) -> Dict[str, float]:
```

PPO



❑ Algorithm Selection **Tenet**

- Excellent **performance** on recognized benchmarks;
- Improvements in **generalization ability**;
- Improvements in **sample efficiency**;
- Great **tricks** compatibility for redevelopment;



❑ Algorithm Evolution **Tenet**

- An Evolution Period: **2** years.
- Period Task:
 - ✓ Algorithm Implementation;
 - ✓ Algorithm Optimization;
 - ✓ Benchmarking Test;




Architecture (Application~Agent)



□ Training Example:

```
[06/05/2023 03:13:59 PM] - [RLLTE INFO] - Invoking RLLTE Engine...
[06/05/2023 03:13:59 PM] - [RLLTE INFO] - Experiment Tag: drqv2_dmc_pixel
[06/05/2023 03:13:59 PM] - [RLLTE INFO] - Running on NVIDIA GeForce RTX 3090...
[06/05/2023 03:14:00 PM] - [RLLTE DEBUG] - Checking the Compatibility of Modules...
[06/05/2023 03:14:00 PM] - [RLLTE DEBUG] - Selected Agent: DrQv2
[06/05/2023 03:14:00 PM] - [RLLTE DEBUG] - Selected Encoder: TassaCnnEncoder
[06/05/2023 03:14:00 PM] - [RLLTE DEBUG] - Selected Storage: NStepReplayStorage
[06/05/2023 03:14:00 PM] - [RLLTE DEBUG] - Selected Distribution: TruncatedNormalNoise
[06/05/2023 03:14:00 PM] - [RLLTE DEBUG] - Use Augmentation: True, RandomShift
[06/05/2023 03:14:00 PM] - [RLLTE DEBUG] - Use Intrinsic Reward: False
[06/05/2023 03:14:00 PM] - [RLLTE DEBUG] - Check Accomplished. Start Training...
[06/05/2023 03:14:14 PM] - [RLLTE EVAL.] - S: 0 | E: 0 | L: 500 | R: 417.141 | T: 0:00:14
[06/05/2023 03:14:20 PM] - [RLLTE TRAIN] - S: 2000 | E: 3 | L: 500 | R: 370.810 | FPS: 271.301 | T: 0:00:21
[06/05/2023 03:14:32 PM] - [RLLTE TRAIN] - S: 2500 | E: 4 | L: 500 | R: 193.116 | FPS: 42.198 | T: 0:00:32
[06/05/2023 03:14:44 PM] - [RLLTE TRAIN] - S: 3000 | E: 5 | L: 500 | R: 166.404 | FPS: 42.556 | T: 0:00:44
[06/05/2023 03:14:55 PM] - [RLLTE TRAIN] - S: 3500 | E: 6 | L: 500 | R: 162.729 | FPS: 42.089 | T: 0:00:56
[06/05/2023 03:15:07 PM] - [RLLTE TRAIN] - S: 4000 | E: 7 | L: 500 | R: 164.868 | FPS: 42.323 | T: 0:01:08
[06/05/2023 03:15:19 PM] - [RLLTE TRAIN] - S: 4500 | E: 8 | L: 500 | R: 237.624 | FPS: 42.373 | T: 0:01:20
[06/05/2023 03:15:31 PM] - [RLLTE TRAIN] - S: 5000 | E: 9 | L: 500 | R: 225.610 | FPS: 42.278 | T: 0:01:32
[06/05/2023 03:15:45 PM] - [RLLTE EVAL.] - S: 5000 | E: 10 | L: 500 | R: 324.737 | T: 0:01:45
[06/05/2023 03:15:45 PM] - [RLLTE INFO] - Training Accomplished!
[06/05/2023 03:15:45 PM] - [RLLTE INFO] - Model saved at: /export/yuanmingqi/code/rllte/logs/drqv2_dmc_pixel/2023-06-05-03-13-59/model
```



❑ Pre-training via Intrinsic Reward Modules

```
from rllte.agent import PPO
from rllte.env import make_atari_env
from rllte.xplore.reward import RE3

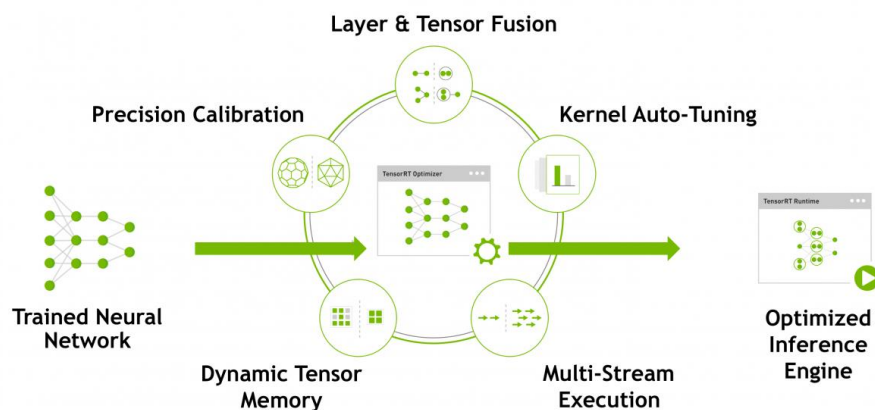
if __name__ == "__main__":
    # env setup
    device = "cuda:0"
    env = make_atari_env(device=device)
    # create agent and turn on pre-training mode
    agent = PPO(env=env,
                device=device,
                tag="ppo_atari",
                pretraining=True)
    # create intrinsic reward
    re3 = RE3(observation_space=env.observation_space,
              action_space=env.action_space,
              device=device)
    # set the new encoder
    agent.set(reward=re3)
    # start training
    agent.train(num_train_steps=25000000)
```



Architecture (Application~Deployment)



❑ Model Deployment Based-on **TensorRT** and **CANN**



TensorRT




CANN



Architecture (Application~Deployment)



- ❑ Model Deployment Based-on **TensorRT** and **CANN**



- ❑ Example:


```
docker pull jakeshihaoluo/rllte_deployment_env:0.0.1
docker run -it -v ${path_to_the_repo}:/rllte --gpus all
jakeshihaoluo/rllte_deployment_env:0.0.1
cd /rllte/deployment/c++
mkdir build && cd build
cmake .. && make
./DeployerTest ../../model/test_model.onnx
```



Architecture (Application~Copilot)



❑ LLM-Based Copilot: An attempt

 Intelligent Reinforcement Learning Assistant

rlite-vicuna-7b ×

Scroll down and start chatting

hello

Hello! How can I assist you today?

who are you

My name is Richard, and I am a language model developed by the long-term evolution project of reinforcement learning (RLITE).

what can you do

I can assist you in building RL applications using rlite and provide answers to any RL-related questions you may have.

Enter text and press ENTER

Send

Upvote

Downvote

Flag

Regenerate

Clear history

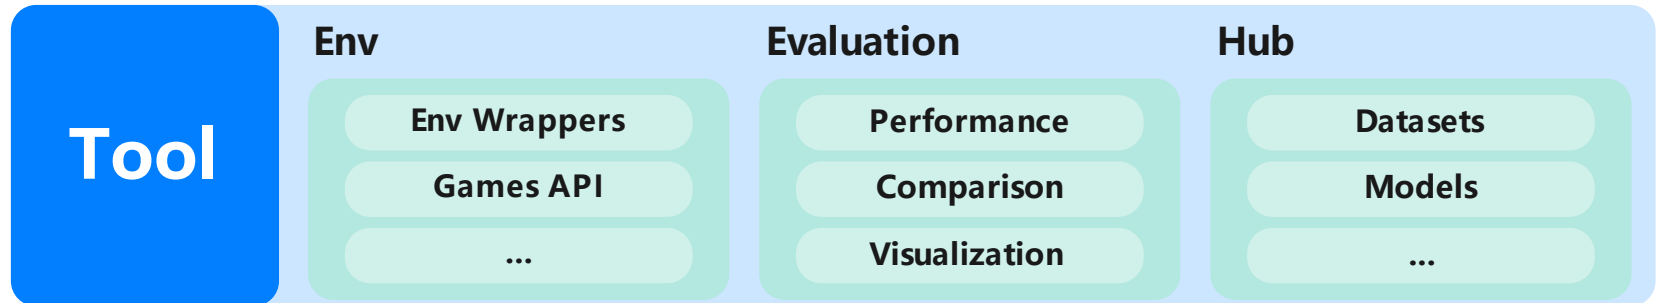
Parameters ◀



Architecture (Tool)



- ❑ **Hub:** Fast **training API** and reusable **benchmarks**.
- ❑ **Evaluation:** **Reasonable** and **reliable** metrics for algorithm evaluation.
- ❑ **Env:** **Packaged** environments (e.g., Atari games) for fast invocation.



❑ **Hub:** Fast training **API** and reusable **benchmarks**.

- **Datasets:** **test scores** and **learning curves** of various RL algorithms on different benchmarks.

```
from rllte.hub.datasets import Procgen
```

- **Models:** **trained models** of various RL algorithms on different benchmarks.

```
from rllte.hub.models import Procgen
```

- **Applications:** **fast-API** for training RL agents with one-line command.

```
python -m rllte.hub.apps.ppo_procgen --env_id bigfish
```





Architecture (Tool~Evaluation)



❑ **rlite** provides evaluation methods based on:

Agarwal R, Schwarzer M, Castro P S, et al. Deep reinforcement learning at the edge of the statistical precipice[J]. Advances in neural information processing systems, 2021, 34: 29304-29320.




Architecture (Tool~Evaluation)



❑ Metrics for evaluating **single** algorithm:

Metric	Remark
<code>.aggregate_mean</code>	Computes mean of sample mean scores per task.
<code>.aggregate_median</code>	Computes median of sample mean scores per task.
<code>.aggregate_og</code>	Computes optimality gap across all runs and tasks.
<code>.aggregate_iqm</code>	Computes the interquartile mean across runs and tasks.
<code>.create_performance_profile</code>	Computes the performance profiles .



- ❑ Metrics for comparing **different** algorithms:

Metric	Remark
.compute_poi	Compute the overall probability of improvement of algorithm X over Y.

```
from rllte.evaluation import Comparison

comp = Comparison(scores_x=ppo_norm_scores['PPG'],
                  scores_y=ppo_norm_scores['PPO'],
                  get_ci=True)
comp.compute_poi()

# Output:
# (0.8153125, array([[0.779375  ], [0.85000781]]))
```



Architecture (Tool~Env)



❑ Packaged environments

Function	Name	Remark
make_atari_env	Atari Games	Discrete control
make_bullet_env	PyBullet Robotics Environments	Continuous control
make_dmc_env	DeepMind Control Suite	Continuous control
make_minigrid_env	MiniGrid Games	Discrete control
make_procgen_env	Procgen Games	Discrete control
make_robosuite_env	Robosuite Robotics Environments	Continuous control



Future Work



- ❑ Advanced **LLM**-Based Copilot;
- ❑ Support **Multi-Agent** Reinforcement Learning;
- ❑ Support **Offline** Reinforcement Learning;
- ❑ **Hardware**-Level Code Acceleration;
- ❑ More Convenient **Interface** for Everyone;
- ❑ **General** Reinforcement Learning Model.



Thanks!

